Agents
Intelligent Systems II

Mário Antunes

2026

Practice Guide: Flappy Bird with Neuro-evolution

In this guide, you will transition from a basic reflex agent to a sophisticated AI capable of learning how to play Flappy Bird using **Neuro-evolution**.

The game architecture is split into two parts:

1. **Backend (Python):** Runs the game physics and logic using asynchronous websockets.
2. **Frontend (HTML5/Canvas):** Visualizes the state sent by the server.

Part 1: Setup

Before writing any code, let's get the game running.

1. **Download the Repository:** Clone the course repository:

   ```
   git clone https://github.com/detiuaveiro/flappy-bird-agent.git
   cd flappy-bird-agent
   ```

2. **Create a Virtual Environment:** It is best practice to isolate your dependencies.

   ```
   python3 -m venv venv
   source venv/bin/activate
   ```

3. **Install Dependencies:** Install the required libraries.

   ```
   pip install -r requirements.txt
   ```

4. **Run the Game:** Start the backend server:

   ```
   python -m src.backend.py --pipes
   ```

Once the server is running, open the `html/play_human.html` file in your web browser. Use the mouse to play the game.

Part 2: The Basic Agent (Reflex Agent)

Your first task is to understand the communication protocol. The server sends the **World State** as a JSON object via websockets.

Task: 1. Review `play_game` method in `play.py`. 2. Locate the message handling loop. You will see a JSON object representing the state. 3. Implement a simple **Rule-Based Agent**. Do not use Machine Learning yet. Write a simple `if` statement.

Example Logic:

```python
# A simple reflex agent
if state['bird_y'] > state['next_pipe_bottom_y'] + 50:
    action = "jump"
else:
    action = "stay"
```

Run this agent and observe. Does it crash? Can it handle different pipe heights?

Part 3: The Intelligent Agent (Neuro-evolution)

Now we move to the core of this course: **Neuro-evolution**.

1. Theoretical Background

Why use Neuro-evolution for Flappy Bird?

- **Continuous World:** The bird's position, velocity, and pipe locations are continuous values (floats).
- **Discrete Action:** The output is binary (Jump or Don't Jump).
- **Partial Observability:** The agent doesn't need to know the entire map; it only needs a "partial view" (distance to the next pipes).

The Solution: Neural Networks + Evolutionary Algorithms

Instead of using Reinforcement Learning (like Q-Learning) which requires a state table (hard for continuous worlds) or Gradient Descent (which requires a dataset of "correct" moves), we use **Neuro-evolution**.

- **The Brain (MLP):** We use a Multi-Layer Perceptron.

- **Input:** Normalized state (Distance to pipe X, Distance to pipe Y, Velocity).

- **Hidden Layers:** Processing units.

- **Output:** A single neuron with a Sigmoid activation. If Output > 0.5 Jump.

- **The Optimization (Evolution):** We do not "train" the network with back-propagation. Instead, we **evolve** the weights of the neural network.

1. **Population:** Generate 50 birds with random weights.
2. **Evaluation:** Let them all play. The "Fitness" is the score (distance traveled).
3. **Selection:** Keep the weights of the birds that flew the furthest.
4. **Crossover & Mutation:** Create a new generation of birds by mixing the weights of the best parents and adding random mutations.

2. Implementation

Open the `train.py` file. The repository is set up to use a Python optimization library (PyBlindOpt; take some time to browse the documentation).

Your Task: Configure the Neuro-evolution loop.

1. **Define the Problem:** Map the Neural Network weights to a 1D vector. The optimization algorithm will try to find the vector of weights that maximizes the game score.
2. **Select an Algorithm:** The library supports several optimization algorithms, such as:

- **DE (Differential Evolution):** Good for global search, maintains diversity.
- **PSO (Particle Swarm Optimization):** Simulates a flock of birds finding food.
- **GWO (Grey Wolf Optimizer):** Mimics the leadership hierarchy of wolves.
- **CS (Cuckoo Search):** Based on the brood parasitism of cuckoo birds.

3. Experiments

Run the training using different configurations and observe the impact on convergence speed (how fast they learn) and stability (do they forget how to fly?).

Experiment A: Algorithm Comparison Run the training for 30 generations with a population of 30 for each algorithm:

- Does **PSO** converge faster than **DE**?
- Does **GWO** find a more stable solution?

Experiment B: Initialization Check the documentation for the optimization library regarding *Initialization Metrics*.

- How does the range of initial random weights (e.g., `[-1, 1]` vs `[-5, 5]`) affect the learning?
- If weights start too large, the sigmoid output might saturate (always 1 or always 0), making the bird constantly jump or never jump.

Experiment C: Advanced Initialization Strategies As discussed in the lectures, standard Random Initialization can result in clusters and gaps in the search space. The library supports advanced strategies to mitigate this.

1. **Opposition-Based Learning (OBL):**

- Enable OBL initialization in your configuration.
- Observation: Compare the "Best Score" in Generation 0 (before any evolution happens) between Random and OBL. Does the OBL population start with a "smarter" agent purely by checking opposite weights?

2. **Sampling Strategies (Sobol / LHS):**

- Switch the sampler to Sobol Sequence or Latin Hypercube Sampling (LHS).
- Analysis: Does covering the search space more evenly reduce the variance between training runs?

3. **OBLESA:**

- Switch to the latest method named OBLESA.
- Does the initial time required for OBLESA to work produces better results (having a bird fly the full limit with less generations)?

Part 4: The Flappy Bird Cup

Friendly Competition: Once you have explored the algorithms, tune your hyperparameters to create an optimal agent.

- **Input features:** Are you feeding the agent `bird_y`? Or `bird_y - pipe_y` (relative distance)? (Hint: Relative distance usually generalizes better).
- **Hidden Layers:** Is a standard `[3 inputs -> 5 hidden -> 1 output]` enough? Or do you need two hidden layers?
- **Population Size:** Larger populations explore better but run slower.
- **Initialization:** Consider using **OBLESA (Opposition-Based with Empty Space Attack)** or **QOBL (Quasi-Opposition)** to get a competitive edge in the early generations.

Goal: Train an agent that can reach **300 point** consistently (fly for 5 minutes).