

# Agents

## Intelligent Systems II

---

Mário Antunes

2026

Universidade de Aveiro

## **Agents and Intelligent Agents**

---

# What is an Agent?

- **Definition:** An agent is anything that can be viewed as perceiving its environment through sensors and acting upon that environment through actuators.
- **The Core Concept:** It is an entity that exercises control over its own actions to achieve goals.
- **Examples:**
  - **Biological:** Humans, animals.
  - **Robotic:** Robotic vacuum cleaners, assembly line arms.
  - **Software:** Web crawlers, NPCs in video games, Trading bots.

# The Agent Interaction Loop i

The fundamental cycle of an agent involves four steps:

1. **Sense**: Gather data from the environment via sensors.
2. **Process/Decide**: The “Black Box” of intelligence processing the input.
3. **Act**: Execute an action via actuators.
4. **Alter**: The action changes the state of the environment.
  - *Repeat*: The agent senses the *new* state.

## The Agent Interaction Loop ii

---

**Analogy:** A pilot flying a plane.

- *Sense:* Reads altimeter and looks out the window.
- *Decide:* "I am too low."
- *Act:* Pulls the yoke back.
- *Alter:* The plane's elevators move, air pressure changes, plane rises.

# Sensors: Perceiving the World

- **Sensors** are the bridge from the physical/virtual world to the agent's internal representation.
- **Percept sequence:** The complete history of everything the agent has perceived.
  - $P = \{p_1, p_2, \dots, p_t\}$
- **Types:**
  - *Passive*: Cameras, Microphones (Receive energy).
  - *Active*: Radar, Lidar, Ultrasound (Emit energy and measure return).
  - *Proprioceptive*: Measuring internal state (Battery level, wheel rotation).

# Actuators: Affecting the World

---

- **Actuators** are the mechanisms that allow the agent to exert force or influence.
- **Mapping:** The agent function maps percepts to actions.
  - $f : P^* \rightarrow A$
- **Examples:**
  - *Physical:* Motors, hydraulic cylinders, speakers, screens.
  - *Virtual:* Sending a network packet, updating a database record, writing a file.

# The Environment

---

The “World” in which the agent operates. It defines the constraints and the physics.

- **State ( $S$ ):** A snapshot of the environment at a specific time  $t$ .
- **Transition:** The environment evolves based on the agent’s actions and its own natural dynamics.
  - $S_{t+1} = T(S_t, A_t)$
- **Key Challenge:** The agent rarely sees the *entire* state (Partial Observability).

# Dimensions of the Environment

To design an intelligent agent, we must classify the environment:

1. **Fully vs. Partially Observable:** Do sensors detect the full state? (Chess vs. Poker).
2. **Deterministic vs. Stochastic:** Is the next state determined purely by current state + action? (Crossword vs. Dice game).
3. **Static vs. Dynamic:** Does the world change while the agent is thinking? (Crossword vs. Tennis).
4. **Discrete vs. Continuous:** (Detailed in next slides).

# Discrete vs. Continuous Worlds

This distinction fundamentally changes the math and algorithms used.

- **Discrete World:**

- Finite number of distinct states.
- Time moves in steps (turns).
- *Math:* Sets, Graphs, State Machines.

- **Continuous World:**

- States are defined by real numbers ( $\mathbb{R}^n$ ).
- Time flows smoothly.
- *Math:* Calculus, Differential Equations, Physics engines.

# Discrete vs. Continuous Actions

---

Just as the world can be granular or smooth, so can the agent's choices.

- **Discrete Actions:**
  - Finite set of choices. e.g.,  $\{Up, Down, Left, Right\}$ .
- **Continuous Actions:**
  - Parameters usually range over continuous values. e.g., Steering angle  $\theta \in [0, 360]$ , Acceleration  $a \in [0, 100]$ .

# World State Definitions

We need more precise terms to describe the complexity of the state space.

Term	Definition	Example
<b>Finite / Tabular</b>	A limited, countable number of states.	Tic-Tac-Toe, Grid World
<b>Infinite / High-Dimensional</b>	Uncountable states. Requires function approximation.	Robot Arm Joint Angles ( $\mathbb{R}^n$ )
<b>Fully Observable</b>	The agent knows the exact state $S_t$ .	Chess, Video Games (HUD)
<b>Partially Observable (POMDP)</b>	The agent only knows a probability distribution.	Poker, Fog of War, Real-world
<b>Structured</b>	State has a known schema/format.	Database Record, JSON
<b>Unstructured / Sub-symbolic</b>	Raw data where features must be extracted.	Pixel data, Audio waveform

# Action Type Definitions

Similarly, actions can be more nuanced than just a simple binary choice.

Term	Definition	Example
<b>Categorical (Discrete)</b>	Selecting one option from a distinct list.	Press {A, B, X, Y}
<b>Parametric (Continuous)</b>	Controlling a value on a sliding scale.	Steer(34.5°)
<b>Hybrid / Parameterized</b>	Choosing a category AND a parameter.	Pass_Ball(Angle=45, Power=0.8)
<b>Deterministic</b>	Action $A$ always leads to State $S'$ .	Move(Right) $\rightarrow$ $(x + 1, y)$
<b>Stochastic</b>	Action $A$ leads to $S'$ with probability $P$ .	Shoot(Target) $\rightarrow$ Hit (80%) / Miss (20%)
<b>Primitive vs. Macro</b>	Atomic actions vs. High-level sequences.	Motor_Voltage(5V) vs. Go_To_Room(B)

# Actions vs. World State: A Taxonomy

Combining these dimensions gives us four distinct types of agent environments.

	Discrete Action	Continuous Action
Discrete World	<b>Chess / Tic-Tac-Toe State:</b> Finite board positions <i>Action:</i> Move piece to square (e4)	<b>Ant Colony Optimization State:</b> Graph nodes (cities) <i>Action:</i> Deposit pheromone intensity $I \in [0, 1]$
Continuous World	<b>Video Games (Platformers) State:</b> ( $x, y, z$ ) coordinates <i>Action:</i> Press Jump Button (0 or 1)	<b>Autonomous Driving / Robotics State:</b> Physics, Friction, Velocity <i>Action:</i> Steering angle $\theta$ , Acceleration $a$

## Case Study 1: Discrete World / Discrete Actions

---

- **Example:** Tic-Tac-Toe or Chess.
- **World:** A 3x3 grid (finite states).
- **Action:** Place 'X' in cell  $(i, j)$ .
- **Why it's "Easy":** We can often enumerate all possibilities (Search Trees).
- **Analogy:** Multiple Choice Question test. You select A, B, C, or D. There is no "Choice A.5".

## Case Study 2: Continuous World / Discrete Actions

- **Example: Modern Video Games** (e.g., a platformer or FPS).
- **World:** Continuous coordinates  $(x, y, z)$  represented by floats. Objects move smoothly.
- **Action:** Buttons on a controller. You press “Jump” or you don’t.
- **Complexity:** The agent must decide *when* to trigger a discrete action in a flowing world.
- **Analogy:** Taking a photo of a race car. The car moves continuously, but the shutter click is a single, discrete event.

## Case Study 3: Continuous World / Continuous Actions

- **Example:** Autonomous Driving or Robotic Arms.
- **World:** The road, physics, friction (Real numbers).
- **Action:**
  - Steering wheel:  $34.5^\circ$
  - Brake pressure: 0.82 Bar
- **Complexity:** Infinite search space. Requires control theory or regression-based learning.
- **Analogy:** Pouring water into a glass. You adjust the tilt continuously to control the flow rate.

# From “Agent” to “Intelligent Agent”

An agent becomes **intelligent** based on how it implements the “Process/Decide” step.

## The Spectrum of Intelligence Sources:

1. **Symbolic/Logic**: Explicit rules and reasoning.
2. **Planning**: Searching for sequences of actions.
3. **Machine Learning**: Learning from data (probabilities).
4. **Deep Learning**: Learning representations.
5. **Agentic AI**: Generative models (LLMs).

## Approach A: Logic and Rule Engines

- **Basis:** First-Order Logic (FOL) and Propositional Logic.
- **Mechanism:** The agent has a Knowledge Base (KB) of facts and rules.
- **Process:**
  1. Sense → Convert to Logical Fact (`dirt(loc_A)`).
  2. Query KB → `?- action(X)`.
  3. Inference → Resolution/Modus Ponens derives `X = suck`.
- **Pros:** Explainable, verifiable.
- **Cons:** Brittle in noisy/unknown environments.

# The Prolog Agent

---

We can use **Prolog** to implement this logic.

- **World Model:** Defined by predicates (at(agent, 1, 1), pit(2, 2)).
- **Reasoning:** prolog  
move(forward) :-  
clear\_ahead, \+ pit\_ahead.  
move(turn\_left) :- wall\_ahead.
- **Logic Revision:** We use Horn Clauses and Resolution to prove which action leads to a goal.

## Approach B: Planning Agents

- **Basis:** Search Algorithms (A\*, BFS) and STRIPS operators.
- **Mechanism:** The agent has a model of "If I do X, Y happens."
- **Process:**
  - *Goal:* Holding(BlockA)
  - *Plan:* Look ahead into the future to find a sequence:  
MoveTo(A) -> PickUp(A).

## Approach C: Machine Learning (RL)

---

- **Basis:** Inductive Inference (Generalizing from examples).
- **Mechanism:** Reinforcement Learning (RL).
- **Feedback Loop:**
  - Agent acts → Environment responds with **Reward** ( $R$ ).
  - Agent updates its policy ( $\pi$ ) to maximize cumulative reward.
- **Analogy:** Dog training. Treat = Reward; Scolding = Punishment. The dog learns “Sit” = “Treat”.

## Approach D: Deep Learning Agents

---

- **The Shift:** Traditional RL struggled with raw inputs (pixels).
- **Deep Q-Networks (DQN):** Use a Neural Network to estimate the value of actions directly from sensors (e.g., video frames).
- **Capability:** Can play Atari games, Drive cars, fold proteins.

## Approach E: Agentic AI (GenAI & LLMs)

---

- **The New Frontier:** Using Large Language Models (LLMs) as the reasoning engine.
- **Mechanism (ReAct Pattern):**
  1. **Thought:** LLM analyzes context. "I need to find the file size."
  2. **Action:** LLM outputs a tool call. `os.stat('file.txt')`
  3. **Observation:** The tool returns output. "2048 bytes".
  4. **Repeat.**
- **Advantage:** Handles unstructured natural language and "common sense" reasoning better than strict logic.

# Summary of Intelligence Architectures

Architecture	Representation	Reasoning Method	Best For
<b>Logic</b>	Symbols/Facts	Theorem Proving	Defined constraints (Tax)
<b>Planning</b>	States/Actions	Search (A*)	Logistics, Robotics
<b>RL/Deep</b>	Vectors/Tensors	Optimization	Motor control, Games
<b>GenAI</b>	Text/Tokens	Probabilistic Generation	Creative/Office tasks

# The Challenge of Continuous Worlds

---

How do we apply Logic or Tabular RL (like Q-Learning) to a continuous world?

- **Problem:** A Q-Table requires a row for every state.
- *Continuous:* Infinite states ( $x = 1.001, x = 1.002\dots$ ).
- *Result:* The table would be infinite. Memory explosion.

## Solution 1: Discretization (Bucketing)

---

Turn the continuous world into a discrete one.

- **Technique:** Divide the continuous range into “buckets” or grid cells.
- **Example:**
  - Real Temp:  $22.4^{\circ}C \rightarrow$  Bucket: “Warm” ( $20 - 25^{\circ}C$ ).
  - Real Dist:  $5.67m \rightarrow$  Grid Cell: (5, 6).
- **Issue: Curse of Dimensionality.** If you have 10 dimensions and divide each into 100 buckets, you need  $100^{10}$  states!

## Solution 2: Function Approximation

---

Instead of remembering every state, learn a **function** that estimates the value.

- **Logic:** Instead of a lookup table  $Q(s, a)$ , use a function  $Q(s, a; \theta) \approx f(s)$ .
- **Linear:**  $Value = w_1 \cdot x + w_2 \cdot y + b$ .
- **Non-Linear:** Neural Networks (Deep Learning).
- **Analogy:** You don't memorize the answer to  $234 \times 567$ ; you learn the *algorithm* of multiplication to handle any numbers.

# Handling Continuous Actions

---

Choosing one action from infinite possibilities.

- **Discretization:** Limit steering to  $\{-10^\circ, 0^\circ, +10^\circ\}$ .  
(Jerky movement).
- **Policy Gradient Methods:**
  - Instead of outputting a specific value, the network outputs the **parameters of a probability distribution** (Mean  $\mu$  and Standard Deviation  $\sigma$ ).
  - Action  $a \sim \mathcal{N}(\mu, \sigma)$ .
- **Actor-Critic:** The “Actor” proposes a continuous action, the “Critic” estimates its value.

# Analogy: The Dart Thrower

---

- **Task:** Throw a dart at the center (Bullseye).
- **World:** Continuous (Physical space).
- **Action:** Continuous (Angle of arm, speed of release).
- **Learning:**
  - *Discrete approach:* Divide the board into squares. "I hit square A4".
  - *Continuous approach:* "I missed by 2cm left." Adjust muscle tension slightly (Gradient Descent).

## Agent Types: Simple Reflex Agent

---

- **Mechanism:** Condition-Action Rules.
- **Logic:** if (Condition) then (Action).
- **Memory:** None. Ignores history.
- **Example:** Thermostat.
  - if (Temp < 20) then (Turn\_On\_Heater).
- **Limitation:** Fails if the environment is partially observable (needs memory to know what is unobserved).

## Agent Types: Model-Based Reflex Agent

---

- **Mechanism:** Keeps internal state to track the world.
- **Logic:** State = Update(OldState, Action, Percept).
- **Example:** Self-driving car waiting at a red light.
  - If a truck blocks the camera, the car remembers “The light was red 2 seconds ago” even if it can’t see it now.
- **Key Component:** “Model of the World” (Physics/Logic).

## Agent Types: Goal-Based Agent

---

- **Mechanism:** Acts to achieve a specific desirable state.
- **Logic:** Search and Planning. "Is action A leading me closer to Goal G?"
- **Difference:** Reflex agents respond to the *past/present*. Goal agents look to the *future*.
- **Example:** GPS Navigation.
  - Goal: "Home".
  - Action: Turn left (because it reduces distance to home, not just because the road turns).

## Agent Types: Utility-Based Agent

- **Mechanism:** Maximizes a “Happiness” function (Utility).
- **Nuance:** Goals are binary (Achieved/Not Achieved).  
Utility is continuous.
- **Example:** Taxi routing.
  - Goal: Get to destination.
  - Utility: Get to destination *quickly AND safely AND save fuel*.
  - Logic: Choose action that maximizes  $E[Utility]$ .

## Recap:

---

We explore specific implementations of these agents:

1. **Representation:** Prolog/Knowledge Graphs (Slide 13).
2. **Communication:** NLP/Grammars (Slide 4).
3. **Learning:**
  - **RL:** Q-Learning/Deep Q-Learning for dynamic worlds (Slide 17).
  - **ILP:** Inductive Logic Programming to learn rules from data (Slide 16).

# Conclusion

Modern AI often combines these approaches (Neuro-Symbolic AI).

- **Sensing:** Deep Learning (Vision).
- **Reasoning:** Logic/Planning (Strategy).
- **Action:** Control Theory (Movement).

**Final Thought:** An intelligent agent is not just an algorithm; it is a system that exists in a loop with its reality.